

HW2 Stat-comp (due Sun, Oct 15th in D2L)

1) Maximum likelihood estimation and inference with the exponential distribution

Recall that the density function of an exponential random variable is

$$f(x_i|\lambda) = \lambda e^{-\lambda x_i}$$

where $x_i \geq 0$ is the random variable, and $\lambda > 0$ is a rate parameter.

The expected value and variance of the random variables are $E[X] = \frac{1}{\lambda}$ and $Var[X] = \frac{1}{\lambda^2}$.

The following code simulates 50 IID draws from an exponential distribution

```
set.seed(195021)
x=rexp(n=50,rate=2)
```

The maximum likelihood estimate of λ has a closed form. Indeed, for a random sample of IID exponentially distributed random variables

$$P(X_1 = x_1, X_2 = x_2, \dots, X_n = x_n | \lambda) = \prod_{i=1}^n \lambda e^{-\lambda x_i} = \lambda^n e^{-\lambda \sum_{i=1}^n x_i} = \lambda^n e^{-\lambda n \bar{x}}$$

Therefore, the log-likelihood is, $l(\lambda|x_1, \dots, x_n) = n \log(\lambda) - \lambda n \bar{x}$.

The derivative of the log-likelihood with respect to λ is

$$\frac{dl}{d\lambda} = \frac{n}{\lambda} - n\bar{x}. \text{ Setting this derivative equal to zero, and solving for } \hat{\lambda} \text{ gives } \hat{\lambda} = \frac{1}{\bar{x}}$$

1.1) Use `optimize()` to estimate λ , compare your estimate with $\frac{1}{\bar{x}}$.

1.2) Use numerical methods to provide an approximate 95% CI for your estimate.

Recall that we can approximate the sampling variance of the maximum likelihood estimate using inverse of the second derivative of the negative log-likelihood evaluated at the ML estimate, also known as the Hessian.

The function `optimize()` does not provide a Hessian. However, you can use the `hessian()` function of the `numDeriv` R-package to obtain a numerical approximation to the second order derivative of the log-Likelihood at the ML estimate. To install this package you can use

```
#install.packages(pkg='numDeriv',repos='https://cran.r-project.org/')
library(numDeriv)
```

To evaluate the Hessian you can use

```
H=hessian(fn=negLogLik,y=y,x=your_ml_estimate)
```

Above, `negLogLik` is a function to evaluate the `-logLik` for the model, `y` is the data and `your_ml_estimate` is the estimate you obtained in 1.1.

To get (an approximation to) the variance of the ML estimate you can use

`VAR=1/H` # if we were estimating a parameter vector we should use `VAR=solve(H)`

Finally, once you have an estimate of the variance you can use estimate $\pm 1.96 \cdot \text{SE}$ to get an approximate 95% CI.

2) CIs for Predictions from Logistic Regression

Recall that in a logistic regression model, the log-odds are parameterized as

$$\log\left[\frac{\theta_i}{1-\theta_i}\right] = \mathbf{x}_i' \boldsymbol{\beta} = \eta_i \quad (1)$$

The sampling variance of $\mathbf{x}_i' \hat{\boldsymbol{\beta}} = \hat{\eta}_i$ is $\text{Var}(\hat{\eta}_i) = \mathbf{x}_i' \mathbf{V} \mathbf{x}_i$, where \mathbf{V} is the (co)variance matrix of the estimated effects; therefore, a SE and an approximate 95%CI for η_i can be obtained using

$$SE(\hat{\eta}_i) = \sqrt{\mathbf{x}_i' \mathbf{V} \mathbf{x}_i} \text{ and } CI : \mathbf{x}_i' \hat{\boldsymbol{\beta}} \pm 1.96 \times SE(\hat{\eta}_i).$$

Because the inverse-logit is a monotonic map, we can then obtain a 95% CI for the predicted probabilities by applying the inverse logit, $\theta_i = \frac{e^{\eta_i}}{1+e^{\eta_i}}$, to the bounds of the CI for the linear predictor.

- Using the gout data set, fit a logistic regression for gout using sex, age, and race as predictors (for this you can use `glm()`, don't forget the link!).
- From the fitted model, derive predictions and SEs for the linear predictor as well as predictions and SEs in the probability scale.

Race	Sex	Age	Predicted LP	95%CI LP	Pred. Prob Scale	95%CI Prob. Scale
White	Male	55				
White	Female	55				
Black	Male	55				
Black	Female	55				

Hint:

- `predict(fm, se.fit=TRUE, newdata=tmp)` returns predictions for the linear predictor, and
- `predict(fm, se.fit=TRUE, newdata=tmp, type="response")` returns predictions and SEs in the probability scale.

Above, `fm` is the fitted logistic regression model, and `newdata=tmp` is a data.frame containing the values of the predictors of the model at which you want to predict.