

HW_2 Solutions

Andriana Manousidaki

2023-10-16

1) Maximum likelihood estimation and inference with the exponential distribution

The density function of an exponential random variable is

$$f(x_i|\lambda) = \lambda e^{-\lambda x_i}$$

where $x_i \geq 0$ is the random variable, and $\lambda > 0$ is a rate parameter.

The expected value and variance of the random variables are $E[X] = \frac{1}{\lambda}$ and $Var[X] = \frac{1}{\lambda^2}$.

The following code simulates 50 IID draws from an exponential distribution

```
set.seed(195021)
x=rexp(n=50,rate=2)
```

The maximum likelihood estimate of λ has a closed form, indeed

$$L(\lambda|x) = \lambda^n e^{-\lambda n\bar{x}}$$

Thus, $l(\lambda|x) = n \log(\lambda) - \lambda n\bar{x}$, therefore

$\frac{dl}{d\lambda} = \frac{n}{\lambda} - n\bar{x}$. Setting this derivative equal to zero, and solving for $\hat{\lambda}$ gives $\hat{\lambda} = \frac{1}{\bar{x}}$

1.1) Use `optimize()` to estimate λ compare your estimate with $\frac{1}{\bar{x}}$.

```
negLogLik=function(y,lambda){
  n=length(x)
  xBar=mean(x)
  logLik=n*log(lambda)-lambda*xBar*n
  return(-logLik)
}
fm=optimize(f=negLogLik,y=x,interval=c(0,100))

fm$minimum
```

```
## [1] 3.247199
```

```
1/mean(x)
```

```
## [1] 3.2472
```

```
fm$objective
```

```
## [1] -8.889658
```

```
negLogLik(y=x,lambda=fm$minimum)
```

```
## [1] -8.889658
```

Or the faster way and a way to reduce calculation errors is to use $P(X=x) = \text{dexp}(x, \text{rate})$

```
negLogLik2=function(y,lambda){  
  logLik=sum(log(dexp(y, rate = lambda)))  
  return(-logLik)  
}  
fm2=optimize(f=negLogLik2,y=x,interval=c(0,100))  
fm2$minimum
```

```
## [1] 3.247199
```

```
1/mean(x)
```

```
## [1] 3.2472
```

```
fm2$objective
```

```
## [1] -8.889658
```

```
negLogLik2(y=x,lambda=fm2$minimum)
```

```
## [1] -8.889658
```

1.2) Use numerical methods to provide an approximate 95% CI for your estimate.

Hint: `optimize()` does not provide a Hessian. However, you can use the `hessian()` function of the `numDeriv` R-package to obtain a numerical approximation to the second order derivative of the logLikelihood at the ML estimate. To install this package you can use

```
#install.packages(pkg='numDeriv',repos='https://cran.r-project.org/')  
library(numDeriv)  
H=hessian(func=negLogLik,y=x,x=fm$minimum)  
H
```

```
##           [,1]  
## [1,] 4.741898
```

```
VAR=1/H # since H is scalar, we just use 1/H  
SE=sqrt(VAR)
```

```
SE
```

```
##           [,1]  
## [1,] 0.4592233
```

```
CI=fm$minimum+c(-1,1)*as.vector(1.96*SE)
round(CI,3)
```

```
## [1] 2.347 4.147
```

2) CIs for Predictions from Logistic Regression

Recall that in a logistic regression model, the log-odds are parameterized as

$$\log\left[\frac{\theta_i}{1-\theta_i}\right] = \mathbf{x}_i'\beta = \eta_i \quad (1)$$

The sampling variance of $\mathbf{x}_i'\beta = \eta_i$ is $Var(\eta_i) = \mathbf{x}_i'\mathbf{V}\mathbf{x}_i$, where \mathbf{V} is the (co)variance matrix of the estimated effects; therefore, a SE and an approximate 95%CI for η_i can be obtained using

$SE(\eta_i) = \sqrt{\mathbf{x}_i'\mathbf{V}\mathbf{x}_i}$ and $CI : \mathbf{x}_i'\hat{\beta} \pm 1.96 \times SE(\eta_i)$.

Because the inverse-logit is a monotonic map, we can then obtain a 95% CI for the predicted probabilities by applying the inverse logit, $\theta_i = \frac{e^{\eta_i}}{1+e^{\eta_i}}$, to the bounds of the CI for the linear predictor.

- Using the gout data set, fit a logistic regression for gout using sex, age, and race as predictors (for this you can use `glm()`, don't forget the link!).
- From the fitted model, and using the formulas presented above, report predictions and 95% CIs in the scale of the linear predictor and in the probability scale.

Race	Sex	Age	Predicted Risk	95%CI
White	Male	55		
White	Female	55		
Black	Male	55		
Black	Female	55		

```
PATH_TO_DATA = 'https://raw.githubusercontent.com/gdlc/STAT_COMP/master/DATA/goutData.txt'
DATA=read.table(PATH_TO_DATA,header=TRUE)
DATA$y=ifelse(is.na(DATA$gout),NA,ifelse(DATA$gout=='Y',1,0))
table(DATA$y,DATA$gout)
```

```
##
##      N   Y
##  0 370   0
##  1   0  30
```

```
fm=glm(y~race+sex+age,data=DATA,family=binomial)
```

Once we fitted the model, we:

- create the incidence matrix (X) for the cases we want to predict,
- evaluate the linear predictor ($\eta = \mathbf{X}\mathbf{b}$), and its SE,
- use the previous results to get a CI for η
- map it, using the inverse-logit, into a CI for the predicted probability

```
## Incidence matrix
X=cbind('int'=1, 'raceW'=c(1,1,0,0), 'sexM'=c(1,0,1,0), 'age'=55)
X
```

```
##      int raceW sexM age
## [1,]   1     1   1  55
## [2,]   1     1   0  55
## [3,]   1     0   1  55
## [4,]   1     0   0  55
```

```
eta=X%%coef(fm)
VCOV.ETA=X%%vcov(fm)%%t(X)
SE.ETA=sqrt(diag(VCOV.ETA))

invLogit=function(eta){
  exp(eta)/(1+exp(eta))
}

LOW=invLogit(eta-1.96*SE.ETA)
UP=invLogit(eta+1.96*SE.ETA)

ANS=data.frame('race'=c('W', 'W', 'B', 'B'), 'sex'=c('M', 'F', 'M', 'F'),
               'age'=55,
               'LP'=eta,
               'LB-LP'=eta-1.96*SE.ETA,
               'UB-LP'=eta+1.96*SE.ETA,
               'prob'=invLogit(eta),
               'LB-prob'=LOW,
               'UB-prob'=UP)

ANS
```

```
##   race sex age      LP      LB.LP      UB.LP      prob      LB.prob      UB.prob
## 1    W  M  55 -3.296722 -4.230460 -2.362984 0.03568382 0.014337159 0.08603923
## 2    W  F  55 -3.745870 -4.705597 -2.786143 0.02307028 0.008963444 0.05807762
## 3    B  M  55 -2.553434 -3.572308 -1.534560 0.07219612 0.027323395 0.17732751
## 4    B  F  55 -3.002582 -3.975647 -2.029517 0.04730937 0.018421435 0.11613854
```

You can also obtain predictions in the probability scale using the `predict()` function, specifying `type=response` and `se.fit=TRUE`. You can then build a 95% CI using prediction $\pm 1.96 \times \text{SE}$. This method can give you predictions outside the $[0,1]$ interval which make no sense. You can then set the lower bound to zero (if the original lower bound was < 0) and the upper bound to 1 (if the original upper bound was > 1).

```
newData=data.frame('race'=c('W', 'W', 'B', 'B'), 'sex'=c('M', 'F', 'M', 'F'), 'age'=55)
TMP=predict(fm, type='response', se.fit=TRUE, newdata = newData)
ANS2=cbind(newData, 'pred.prob'=TMP$fit,
            'SE'=TMP$se.fit,
            'LB'=TMP$fit-1.96*TMP$se.fit,
            'UB'=TMP$fit+1.96*TMP$se.fit)

# Fixing possible bounds below 0 or above 1
ANS2$LB=ifelse(ANS2$LB<0, 0, ANS2$LB)
```

```
ANS2$UB=ifelse(ANS2$UB>1,1,ANS2$UB)
```

```
ANS2
```

##	race	sex	age	pred.prob	SE	LB	UB
## 1	W	M	55	0.03568382	0.01639304	0.003553453	0.06781418
## 2	W	F	55	0.02307028	0.01103590	0.001439908	0.04470064
## 3	B	M	55	0.07219612	0.03482046	0.003948008	0.14044423
## 4	B	F	55	0.04730937	0.02237613	0.003452156	0.09116659